



## Data Loading & Initial Overview

### Load Dataset

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")

from google.colab import files
uploaded = files.upload()
df = pd.read_excel(list(uploaded.keys())[0])
```

Upload widget is only available when the cell has been executed

in the current browser session. Please rerun this cell to enable.

Saving AirQualityUCI.xlsx to AirQualityUCI (1).xlsx

### Initial Overview

*Purpose:*

Understand structure

Identify missing values

Identify data types

```
In [ ]: df.head()
df.shape
df.info()
df.describe()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9357 entries, 0 to 9356
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date              9357 non-null    datetime64[ns]
 1   Time              9357 non-null    object  
 2   CO(GT)            9357 non-null    float64 
 3   PT08.S1(CO)       9357 non-null    float64 
 4   NMHC(GT)          9357 non-null    int64   
 5   C6H6(GT)          9357 non-null    float64 
 6   PT08.S2(NMHC)     9357 non-null    float64 
 7   NOx(GT)           9357 non-null    float64 
 8   PT08.S3(NOx)      9357 non-null    float64 
 9   NO2(GT)           9357 non-null    float64 
 10  PT08.S4(NO2)      9357 non-null    float64 
 11  PT08.S5(03)       9357 non-null    float64 
 12  T                 9357 non-null    float64 
 13  RH                9357 non-null    float64 
 14  AH                9357 non-null    float64 

dtypes: datetime64[ns](1), float64(12), int64(1), object(1)
memory usage: 1.1+ MB

```

Out[ ]:		Date	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)
	<b>count</b>	9357	9357.000000	9357.000000	9357.000000	9357.000000
	<b>mean</b>	2004-09-21 04:30:05.193972480	-34.207524	1048.869652	-159.090093	1.865576
	<b>min</b>	2004-03-10 00:00:00	-200.000000	-200.000000	-200.000000	-200.000000
	<b>25%</b>	2004-06-16 00:00:00	0.600000	921.000000	-200.000000	4.004958
	<b>50%</b>	2004-09-21 00:00:00	1.500000	1052.500000	-200.000000	7.886653
	<b>75%</b>	2004-12-28 00:00:00	2.600000	1221.250000	-200.000000	13.636091
	<b>max</b>	2005-04-04 00:00:00	11.900000	2039.750000	1189.000000	63.741476
	<b>std</b>	NaN	77.657170	329.817015	139.789093	41.380154

## Data Cleaning & Pre-processing

### Handling Missing Values

```
In [ ]: df.replace(-200, np.nan, inplace=True)
```

### Correcting Date & Time Data Types

```
In [ ]: df['Date'] = pd.to_datetime(df['Date'], dayfirst=True)

def fix_time(x):
    if isinstance(x, str):
        return pd.to_datetime(x).time()
    return x

df['Time'] = df['Time'].apply(fix_time)
```

## Creating Derived Columns

```
In [ ]: df['Hour'] = df['Time'].apply(lambda x: x.hour)
df['Month'] = df['Date'].dt.month
df['Year'] = df['Date'].dt.year
```

## Removing Duplicates

```
In [ ]: df.drop_duplicates(inplace=True)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9357 entries, 0 to 9356
Data columns (total 18 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Date              9357 non-null   datetime64[ns]
 1   Time              9357 non-null   object 
 2   CO(GT)            7674 non-null   float64
 3   PT08.S1(CO)       8991 non-null   float64
 4   NMHC(GT)          914 non-null    float64
 5   C6H6(GT)          8991 non-null   float64
 6   PT08.S2(NMHC)     8991 non-null   float64
 7   NOx(GT)           7718 non-null   float64
 8   PT08.S3(NOx)      8991 non-null   float64
 9   NO2(GT)           7715 non-null   float64
 10  PT08.S4(NO2)      8991 non-null   float64
 11  PT08.S5(03)       8991 non-null   float64
 12  T                 8991 non-null   float64
 13  RH                8991 non-null   float64
 14  AH                8991 non-null   float64
 15  Hour              9357 non-null   int64  
 16  Month             9357 non-null   int32  
 17  Year              9357 non-null   int32  
dtypes: datetime64[ns](1), float64(13), int32(2), int64(1), object(1)
memory usage: 1.2+ MB
```

## MISSING VALUES VERIFICATION

```
In [ ]: df.isna().sum()
```

```
Out[ ]:          0
                Date    0
                Time    0
                CO(GT) 1683
                PT08.S1(CO) 366
                NMHC(GT) 8443
                C6H6(GT) 366
                PT08.S2(NMHC) 366
                NOx(GT) 1639
                PT08.S3(NOx) 366
                NO2(GT) 1642
                PT08.S4(NO2) 366
                PT08.S5(O3) 366
                T      366
                RH     366
                AH     366
                Hour   0
                Month  0
                Year   0
```

**dtype:** int64

SHAPE BEFORE vs AFTER CLEANING

```
In [ ]: print("Final dataset shape:", df.shape)
```

```
Final dataset shape: (9357, 18)
```


The dataset was cleaned by handling missing values, correcting data types, and removing duplicates, resulting in a structured dataset suitable for exploratory analysis.

## TASK 4: Exploratory Data Analysis (EDA) & Visualizations

### 1. Air Quality Analysis

#### 1.1 Distribution of CO Concentration - Univariate


```
In [ ]: sns.histplot(df['CO(GT)'], bins=30)
plt.title('Distribution of CO Concentration')
plt.show()
```



Purpose: Assess pollution severity and variability.


#### 1.2.1 Hourly Pollution Pattern

```
In [ ]: df.groupby('Hour')['CO(GT)'].mean().plot()
plt.title('Hourly CO Variation')
plt.show()
```



### 1.3 Monthly Pollution Trend


```
In [ ]: df.groupby('Month')['NO2(GT)'].mean().plot()  
plt.title('Monthly NO2 Trend')  
plt.show()
```



#### 1.4 Bivariate Analysis – Pollutant Relationships

```
In [ ]: numeric_df = df.select_dtypes(include=[np.number])

sns.heatmap(numeric_df.corr(), cmap='coolwarm')
plt.title('Correlation Between Air Pollutants')
plt.show()
```




Bivariate analysis of pollutant relationships examines how two pollutants vary together to identify common emission sources and understand air-quality behavior.

### 1.5 Daily Air Quality Trend (Time-Series) Purpose

Shows overall pollution evolution over time (real air quality, not sensors). Identifies prolonged high-pollution periods and long-term trends.

```
In [ ]: daily_co = df.groupby('Date')['CO(GT)'].mean()

plt.figure(figsize=(12,4))
daily_co.plot()
plt.title('Daily Average CO Concentration')
plt.xlabel('Date')
plt.ylabel('CO (GT)')
plt.show()
```




1.6 Diurnal air-quality behavior across pollutants. Shows traffic-related pollution peaks clearly.

```
In [ ]: hourly_pollutants = df.groupby('Hour')[['CO(GT)', 'NO2(GT)', 'NOx(GT)']].mean()

plt.figure(figsize=(10,5))
hourly_pollutants.plot()
plt.title('Hourly Variation of Major Pollutants')
plt.xlabel('Hour')
plt.ylabel('Concentration')
plt.show()
```

<Figure size 1000x500 with 0 Axes>




```
In [ ]: pollutant_cols = ['CO(GT)', 'NO2(GT)', 'NOx(GT)', 'C6H6(GT)']
sns.pairplot(df[pollutant_cols])
plt.suptitle('Pairwise Relationships Between Pollutants', y=1.02)
plt.show()
```

## 1.7 Relative Contribution of Pollutants

```
In [ ]: mean_pollutants = df[['CO(GT)', 'NO2(GT)', 'NOx(GT)', 'C6H6(GT)']].mean()

plt.figure()
mean_pollutants.plot(kind='pie', autopct='%1.1f%%')
plt.title('Relative Contribution of Pollutants')
plt.ylabel('')
plt.show()
```

## Relative Contribution of Pollutants



## 2. Multi Sensor Detectors - Analysis

### 2.1 Performance

```
In [ ]: performance_corr = pd.DataFrame({
    'CO Sensor (S1)': abs(df[['CO(GT)', 'PT08.S1(CO)']].corr().iloc[0,1]) * 100,
    'NMHC Sensor (S2)': abs(df[['NMHC(GT)', 'PT08.S2(NMHC)']].corr().iloc[0,1]),
    'NOx Sensor (S3)': abs(df[['NOx(GT)', 'PT08.S3(NOx)']].corr().iloc[0,1]) * 100,
    'NO2 Sensor (S4)': abs(df[['NO2(GT)', 'PT08.S4(NO2)']].corr().iloc[0,1]) * 100
}, index=['Performance %']).T


performance_corr
```

```
Out[ ]:
```

Performance %	
<b>CO Sensor (S1)</b>	87.929166
<b>NMHC Sensor (S2)</b>	87.769055
<b>NOx Sensor (S3)</b>	65.569039
<b>NO2 Sensor (S4)</b>	15.767751

```
In [ ]: plt.figure(figsize=(8,4))
final_table['Performance (%)'].plot(kind='bar')
plt.ylabel('Performance (%)')
plt.title('Performance Comparison of Low-Cost Sensors')
```

```
plt.xticks(rotation=45, ha='right')
plt.show()
```



## 2.2 Error

```
In [ ]: from scipy.stats import zscore
import numpy as np

def normalized_mae(gt, sensor):
    # Align and drop NaNs together
    aligned = pd.concat([gt, sensor], axis=1).dropna()

    gt_norm = zscore(aligned.iloc[:, 0])
    sensor_norm = zscore(aligned.iloc[:, 1])

    return np.mean(np.abs(sensor_norm - gt_norm))
```

```
In [ ]: norm_error = {
    'CO Sensor (S1)': normalized_mae(df['CO(GT)'], df['PT08.S1(CO)']),
    'NMHC Sensor (S2)': normalized_mae(df['NMHC(GT)'], df['PT08.S2(NMHC)']),
    'NOx Sensor (S3)': normalized_mae(df['NOx(GT)'], df['PT08.S3(Nox)']),
    'NO2 Sensor (S4)': normalized_mae(df['NO2(GT)'], df['PT08.S4(NO2)'])
}


norm_error_df = pd.DataFrame.from_dict(
    norm_error, orient='index', columns=['Normalized Error']
)
```

```
norm_error_df
```

Out[ ]:

Normalized Error	
<b>CO Sensor (S1)</b>	0.362357
<b>NMHC Sensor (S2)</b>	0.383716
<b>NOx Sensor (S3)</b>	1.389986
<b>NO2 Sensor (S4)</b>	1.070079

```
In [ ]: plt.figure(figsize=(8,4))
final_table['Error'].plot(kind='bar')
plt.ylabel('Error (Normalized)')
plt.title('Error Comparison of Low-Cost Sensors')
plt.xticks(rotation=45, ha='right')
plt.show()
```



```
In [ ]: min_err = norm_error_df['Normalized Error'].min()
max_err = norm_error_df['Normalized Error'].max()

norm_error_df['Scaled Error %'] = (
    (norm_error_df['Normalized Error'] - min_err) /
    (max_err - min_err)
) * 100
```

```
In [ ]: performance = corr_df.iloc[:, 0]
error = norm_error_df['Scaled Error %']

In [ ]: final_table = pd.concat([performance, error], axis=1)

In [ ]: final_table.columns = ['Performance (%)', 'Error (%)']

In [ ]: corr_df.index = [
    'CO Sensor (S1)',
    'NMHC Sensor (S2)',
    'NOx Sensor (S3)',
    'NO2 Sensor (S4)'
]
```

### 2.3 Performance V/S Error

```
In [ ]: final_table = pd.concat(
    [
        corr_df.iloc[:, 0],           # Performance
        norm_error_df.iloc[:, 0]      # Error (Normalized Error)
    ],
    axis=1
)

final_table.columns = ['Performance (%)', 'Error']

final_table
```

Out[ ]:


	Performance (%)	Error
<b>CO Sensor (S1)</b>	87.929166	0.362357
<b>NMHC Sensor (S2)</b>	87.769055	0.383716
<b>NOx Sensor (S3)</b>	65.569039	1.389986
<b>NO2 Sensor (S4)</b>	15.767751	1.070079

```
In [ ]: plt.figure(figsize=(6,5))
plt.scatter(final_table['Error'], final_table['Performance (%)'])

for sensor in final_table.index:
    plt.text(
        final_table.loc[sensor, 'Error'],
        final_table.loc[sensor, 'Performance (%)'],
        sensor,
        fontsize=9
    )

plt.xlabel('Error')
plt.ylabel('Performance (%)')
plt.title('Performance vs Error of Low-Cost Sensors')
```

```
plt.show()
```



```
In [ ]: final_table['Performance Rank'] = final_table['Performance (%]').rank(ascending=False)
final_table['Error Rank'] = final_table['Error'].rank(ascending=True)
final_table['Total Score'] = final_table['Performance Rank'] + final_table['Er
```

```
In [ ]: # Sort by Total Score (best first)
final_table = final_table.sort_values('Total Score')

# Assign grades
grades = ['Best', 'Better', 'Good', 'Average']
final_table['Grade'] = grades[:len(final_table)]

final_table
```

Out[ ]:

	Performance (%)	Error	Performance Rank	Error Rank	Total Score	Best Sensor	Grade
CO Sensor (S1)	87.929166	0.362357	1.0	1.0	2.0	Yes	Best
NMHC Sensor (S2)	87.769055	0.383716	2.0	2.0	4.0	No	Better
NOx Sensor (S3)	65.569039	1.389986	3.0	4.0	7.0	No	Good
NO2 Sensor (S4)	15.767751	1.070079	4.0	3.0	7.0	No	Average

In [ ]: `final_table = final_table[['Performance (%)', 'Error', 'Grade']]  
final_table`

Out[ ]:

	Performance (%)	Error	Grade
CO Sensor (S1)	87.929166	0.362357	Best
NMHC Sensor (S2)	87.769055	0.383716	Better
NOx Sensor (S3)	65.569039	1.389986	Good
NO2 Sensor (S4)	15.767751	1.070079	Average

Table compares the performance and error characteristics of low-cost sensors relative to reference analyzers. Performance was evaluated using correlation-based trend accuracy, while error was quantified using normalized and scaled error metrics. The CO sensor (PT08.S1) exhibits the highest overall performance and the lowest relative error, indicating superior reliability among the evaluated low-cost sensors.

## Insights

1. **CO Concentration Distribution (Histogram)** Most CO values are low, with few high peaks, showing occasional pollution events.
2. **Hourly CO Variation (Diurnal Plot)** CO levels peak in the morning and evening due to traffic activity.
3. **Monthly NO<sub>2</sub> Trend** NO<sub>2</sub> concentration is higher in winter and lower in summer, showing seasonal effects.

4. **Correlation Heatmap of Pollutants** CO, NOx, and NO<sub>2</sub> are strongly correlated, indicating common emission sources.
5. **Daily Average CO Time Series** CO levels vary gradually over time, indicating sensor drift rather than sudden faults.
6. **Hourly Variation of Multiple Pollutants** All pollutants show similar daily patterns, confirming traffic influence.
7. **Scatter Plot Between CO and NOx** CO and NOx increase together, showing a strong linear relationship.
8. **Scatter Plot Between NOx and NO<sub>2</sub>** NO<sub>2</sub> closely follows NOx levels, confirming combustion-related sources.
9. **Pollutant Contribution (Pie Chart)** CO and NOx contribute the largest share of air pollution.
10. **Sensor Performance Comparison (Bar Chart)** Sensor accuracy varies, showing that some sensors perform better than others.
11. **Sensor Error Comparison** Sensors with lower accuracy show higher measurement errors.
12. **Performance vs Error Plot** Higher sensor performance corresponds to lower error values.
13. CO Sensor performs well and having less error rate among all other sensors.

In [ ]: